

Documentation for app.py - Dental Clinic Receptionist AI Assistant

Overview

The app.py file implements an AI-powered dental clinic receptionist that interacts with users via voice. It listens to user speech, transcribes it, generates a response, and converts the response to audio. The application uses free, open-source tools to ensure no paid services are required. It is designed to be lightweight and compatible with Docker for deployment.

Features

- **Speech-to-Text:** Captures user speech using a microphone and transcribes it with Google's Speech-to-Text API via the speech_recognition library.
- **Text Generation:** Generates responses using Hugging Face's distilgpt2 model, a lightweight transformer model.
- **Text-to-Speech:** Converts text responses to audio using gTTS (Google Text-to-Speech).
- **Audio Playback:** Plays audio responses using pygame, which is Docker-compatible.
- **Error Handling:** Robustly handles transcription, generation, and playback errors, with graceful shutdown on interrupt.
- **Free Tools:** Uses only free libraries and services, requiring no API keys or paid subscriptions.

Dependencies

- Python 3.10 or 3.11 (Python 3.13.3 may have compatibility issues with PyTorch)
- Libraries:
 - speechrecognition>=3.10.0: For speech-to-text transcription.
 - pyaudio>=0.2.14: For microphone input.
 - transformers>=4.35.0: For text generation with distilgpt2.
 - gTTS>=2.3.2: For text-to-speech conversion.
 - pygame>=2.5.2: For audio playback.
 - python-dotenv>=1.0.0: For loading environment variables (optional, as no API keys are needed).
 - torch>=2.4.1: Required by transformers for model execution.

Setup Instructions

1. Install Python

- **Recommended Version:** Python 3.10 or 3.11 (due to compatibility with PyTorch and transformers).
- **Windows:**
 - Download and install Python 3.11 from python.org.
 - Ensure pip and python are added to your PATH during installation.
- **Verify Installation:**
 - python --version
 - pip --version

2. Create a Virtual Environment

- Create a virtual environment to isolate dependencies:
- python -m venv venv
- Activate the virtual environment:
- .\venv\Scripts\Activate.ps1

3. Install Dependencies

- Create a requirements.txt file with the following content:
- speechrecognition>=3.10.0
- pyaudio>=0.2.14
- transformers>=4.35.0
- gTTS>=2.3.2
- pygame>=2.5.2
- python-dotenv>=1.0.0
- torch>=2.4.1
- Install dependencies:
- pip install -r requirements.txt
- **Note:** If torch installation fails with Python 3.13.3, switch to Python 3.11 and recreate the virtual environment.

4. Create .env File (Optional)

- No API keys are required, but the code checks for a .env file. Create an empty file or skip this step:

- # .env
- # No API keys required

5. Test Microphone and Audio

- Ensure your microphone and speakers are working (test with a tool like Audacity).
- Verify internet connectivity for Google Speech-to-Text.

Code Explanation

Imports

`import os`

`from dotenv import load_dotenv`

`import speech_recognition as sr`

`from transformers import pipeline`

`from gtts import gTTS`

`import pygame`

`import tempfile`

`import time`

- `os`: For file operations (e.g., deleting temporary files).
- `dotenv`: Loads environment variables (optional).
- `speech_recognition`: Captures and transcribes audio.
- `transformers`: Generates text responses using `distilgpt2`.
- `gTTS`: Converts text to speech.
- `pygame`: Plays audio files.
- `tempfile`: Creates temporary files for audio.
- `time`: Adds delays for error handling.

Environment Setup

`load_dotenv()`

- Loads environment variables from `.env`. No variables are required, so this is a no-op if the file is absent.

AI_Assistant Class

The core class that manages the assistant's functionality.

`__init__`

```
def __init__(self):
```

```
    self.recognizer = sr.Recognizer()
```

```
    self.microphone = sr.Microphone()
```

```
    self.text_generator = pipeline("text-generation", model="distilgpt2")
```

```
    pygame.mixer.init()
```

```
    self.full_transcript = [
```

```
        {"role": "system", "content": "You are a receptionist at a dental clinic. Be resourceful and efficient."},
```

```
    ]
```

- Initializes the speech recognizer and microphone.
- Sets up the distilgpt2 model for text generation.
- Initializes pygame for audio playback.
- Creates a conversation history (full_transcript) with a system prompt defining the assistant's role.

`start_transcription`

```
def start_transcription(self):
```

```
    print("Listening... Speak now.")
```

```
    with self.microphone as source:
```

```
        self.recognizer.adjust_for_ambient_noise(source)
```

```
        while True:
```

```
            try:
```

```
                audio = self.recognizer.listen(source, timeout=5, phrase_time_limit=10)
```

```
            try:
```

```
                transcript = self.recognizer.recognize_google(audio)
```

```
            if transcript.strip():
```

```
                print(f"\nPatient: {transcript}\n")
```

```
                self.generate_ai_response(transcript)
```

```

        else:
            print("No speech detected, continuing to listen...")
    except sr.UnknownValueError:
        print("Could not understand audio, continuing to listen...")
    except sr.RequestError as e:
        print(f"Google API error: {e}, continuing to listen...")
except KeyboardInterrupt:
    print("\nStopping transcription...")
    break
except Exception as e:
    print(f"Error during transcription: {e}")
    time.sleep(1)

```

- Listens for audio input in a loop.
- Adjusts for ambient noise to improve transcription accuracy.
- Captures audio with a 5-second timeout and 10-second phrase limit.
- Transcribes audio using Google's Speech-to-Text API.
- Handles errors (unintelligible audio, API failures, or other exceptions).
- Calls `generate_ai_response` for valid transcriptions.
- Exits on KeyboardInterrupt (Ctrl+C).

generate_ai_response

```

def generate_ai_response(self, transcript):
    self.full_transcript.append({"role": "user", "content": transcript})

    prompt = f"System: You are a receptionist at a dental clinic. Be resourceful and efficient.\nPatient: {transcript}\nReceptionist:"

    try:
        response = self.text_generator(prompt, max_length=100, num_return_sequences=1,
truncation=True)[0]["generated_text"]

        ai_response = response.split("Receptionist:")[1].strip()

        self.full_transcript.append({"role": "assistant", "content": ai_response})

```

```
self.generate_audio(ai_response)
```

```
except Exception as e:
```

```
    print(f"Error generating AI response: {e}")
```

- Appends the user's transcript to the conversation history.
- Creates a prompt for distilgpt2 with the system role and user input.
- Generates a response (max 100 tokens) and extracts the assistant's part.
- Appends the response to the conversation history.
- Calls generate_audio to convert the response to speech.
- Handles errors during generation.

generate_audio

```
def generate_audio(self, text):
```

```
    print(f"\nAI Receptionist: {text}")
```

```
    try:
```

```
        tts = gTTS(text=text, lang="en")
```

```
        with tempfile.NamedTemporaryFile(delete=False, suffix=".mp3") as temp_audio_file:
```

```
            tts.save(temp_audio_file.name)
```

```
            temp_audio_file_path = temp_audio_file.name
```

```
        pygame.mixer.music.load(temp_audio_file_path)
```

```
        pygame.mixer.music.play()
```

```
        while pygame.mixer.music.get_busy():
```

```
            pygame.time.Clock().tick(10)
```

```
        pygame.mixer.music.unload()
```

```
        os.remove(temp_audio_file_path)
```

```
    except Exception as e:
```

```
        print(f"Error generating or playing audio: {e}")
```

- Prints the assistant's response.
- Converts text to speech using gTTS and saves it as a temporary MP3 file.
- Loads and plays the audio using pygame.

- Waits for playback to finish.
- Unloads the audio and deletes the temporary file.
- Handles errors during audio generation or playback.

__del__

```
def __del__(self):
```

```
    pygame.mixer.quit()
```

- Cleans up pygame resources when the object is destroyed.

Main Execution

```
if __name__ == "__main__":
```

```
    greeting = "Thank you for calling Vancouver dental clinic. My name is Sandy, how may I assist you?"
```

```
    try:
```

```
        ai_assistant = AI_Assistant()
```

```
        ai_assistant.generate_audio(greeting)
```

```
        ai_assistant.start_transcription()
```

```
    except KeyboardInterrupt:
```

```
        print("\nShutting down...")
```

- Defines a greeting message.
- Initializes the assistant, plays the greeting, and starts transcription.
- Handles KeyboardInterrupt for clean shutdown.

Running the Application

Locally

1. Set Up Environment:

- Ensure Python 3.10 or 3.11 is installed.
- Create and activate a virtual environment:
- `python -m venv venv`
- `.\venv\Scripts\Activate.ps1`

2. Install Dependencies:

- Create requirements.txt (see above).
- Install dependencies:
- `pip install -r requirements.txt`
- If torch installation fails, try:
- `pip install torch==2.4.1`
- If issues persist, switch to Python 3.11 and recreate the virtual environment.

3. Run the Application:

- Ensure your microphone and speakers are working.
- Run the script:
- `python app.py`
- The assistant will play a greeting, then listen for your speech. Speak clearly, and it will transcribe, respond, and play the response.

4. Stop the Application:

- Press Ctrl+C to stop the application.

In Docker

1. Create Dockerfile:

2. # Use a Python base image
3. FROM python:3.10-slim
- 4.
5. # Set working directory
6. WORKDIR /app
- 7.
8. # Install system dependencies for audio (pygame, pyaudio)
9. RUN apt-get update && apt-get install -y \
10. libsndfile1 \
11. libportaudio2 \
12. portaudio19-dev \
13. && rm -rf /var/lib/apt/lists/*

- 14.
15. # Copy requirements file
16. COPY requirements.txt .
- 17.
18. # Install Python dependencies
19. RUN pip install --no-cache-dir -r requirements.txt
- 20.
21. # Copy the application code
22. COPY app.py .
- 23.
24. # Command to run the application
25. CMD ["python", "app.py"]
26. **Create requirements.txt** (see above).
27. **Build the Docker Image:**
28. docker build -t dental-ai-assistant .
29. **Run the Docker Container:**
 - Ensure audio and microphone access:
 - docker run --rm -it --device=/dev/snd --network=host dental-ai-assistant
 - On Windows, you may need to run in WSL2 or adjust audio device access:
 - # In WSL2
 - wsl docker run --rm -it --device=/dev/snd --network=host dental-ai-assistant
30. **Stop the Container:**
 - Press Ctrl+C to stop the container.

Troubleshooting

Installation Issues

- **PyTorch Installation Fails:**
 - If pip install torch fails with Python 3.13.3, switch to Python 3.11:
 - # Install Python 3.11 via python.org or pyenv

- `python -m venv venv`
- `.\venv\Scripts\Activate.ps1`
- `pip install -r requirements.txt`
- Try a specific PyTorch version:
- `pip install torch==2.4.1`
- Check [PyTorch's installation guide](#).
- **Dependency Conflicts:**
 - Ensure all dependencies are compatible by using the versions in `requirements.txt`.
 - If conflicts occur, recreate the virtual environment and reinstall.

Runtime Issues

- **No Transcription:**
 - Verify your microphone is working (test with Audacity).
 - Ensure `pyaudio` is installed:
 - `pip install pyaudio`
 - Check internet connectivity for Google Speech-to-Text.
 - If Google API fails repeatedly, it may be due to usage limits. Try again later.
- **No Audio Playback:**
 - Ensure `pygame` is installed and speakers are working.
 - In Docker, verify `--device=/dev/snd` is passed:
 - `docker run --rm -it --device=/dev/snd --network=host dental-ai-assistant`
- **Poor AI Responses:**
 - `distilgpt2` is lightweight and may produce less coherent responses. Switch to `gpt2-medium` for better quality (requires more memory):
 - `self.text_generator = pipeline("text-generation", model="gpt2-medium")`
 - Fine-tuning the model is possible but requires additional setup.
- **Docker Issues:**
 - If audio or microphone access fails, try `--network=host` or additional device permissions.

- Ensure libportaudio2 and portaudio19-dev are installed in the container.

Limitations

- **Google Speech-to-Text:** Requires internet access and has usage limits. Responses may fail if limits are exceeded.
- **distilgpt2:** Limited conversational ability compared to advanced models like GPT-3.5. Suitable for simple receptionist tasks.
- **Real-Time Performance:** Slight delays may occur due to Google API latency or model inference.
- **Python Compatibility:** Python 3.13.3 may cause issues with PyTorch or transformers. Use Python 3.10 or 3.11 for stability.

Example Usage

1. Run the application:
2. `python app.py`
3. Hear the greeting: "Thank you for calling Vancouver dental clinic. My name is Sandy, how may I assist you?"
4. Speak a query (e.g., "I need to book an appointment").
5. The assistant transcribes your speech, generates a response (e.g., "Sure, I can help with that. When would you like to come in?"), and plays it back.
6. Press Ctrl+C to stop.

Future Improvements

- **Better Transcription:** Integrate a local speech-to-text model (e.g., Vosk) for offline support.
- **Improved Responses:** Use a larger model like gpt2-medium or fine-tune distilgpt2 for better coherence.
- **Multi-Turn Conversation:** Enhance full_transcript to maintain context over multiple interactions.
- **Docker Audio:** Streamline audio configuration for Windows and non-WSL2 environments.